

Topic Proposal: Introducing and Comparing NoSQL Database Management Systems

Jiarong Ye

September 7, 2018

Abstract

With the exponentially generated data from varied sources like mobile devices, mechanical sensors, financial transactions, satellite imaging, the task of data managing gets much more challenging day after day. And these collected raw data are in a wide range of forms, most are unstructured, like images, videos, audios, etc. In the whole data processing pipeline, starting from data integrating and storing, it's already a quite essential part. Since the consistency, accessibility, velocity of data inserting and retrieving might make a huge difference to the following data manipulation tasks, especially in fields like finance, social networks, etc. Thus making an optimal choice of the right database for different data format, different data volume, different business focus like whether velocity or consistency is more important for the client's perspective is becoming more essential than ever.

Widely used Database Management Systems are Relational DBMS such as SQL (MySQL, SQL Server), and NoSQL(Not Only SQL). Why the relational DBMS is limited is that its strong focus on consistency and reliability and ACID guaranteed (ACID: Atomicity, Consistency, Isolation and Durability) also make it difficult for scaling, which is one of the many reasons why NoSQL is considered better than SQL. To reach a balance between availability and consistency, NoSQL achieve more optimal level of consistency by slightly sacrificing its availability.

Goal

The proposed goal of the paper is to give a precise introduction of the NoSQL DBMS and provide a comparison of some of the most frequently used ones in the industries nowadays, Cassandra, HBase and MongoDB, mainly from 4 aspects: their mechanism, important features, advantages, existing problems.

Proposing Outline

NoSQL DBMS (very rough example, more details will be concluded) There are mainly 4 types of NoSQL databases:

- Key / Value Store. It the simplest NoSQL databases. Every item in the database is stored as a key together with its value.
 - e.g. Redis, Riak

- Column Store. Instead of storing rows, it stores columns of data as entries so it's preferred for data manipulations on larger datasets.
 - e.g. Cassandra, HBase
- Document store. It stores data with a more complicated data structure known as documents, which could contain key-value pairs, or key-array pairs, or nested documents.
 - e.g. MongoDB, Couchbase
- Graph databases. It's normally used for networks data storing, such as social networks.
 - e.g. OrientDB, Neo4J

Comparison From the varied industrial level used NoSQL DBMS, we mainly discuss the following 3 and would make comprehensive comparisons among these DBMS:

- Cassandra
- MongoDB
- Hbase

on factors listed below:

- mechanism
 - Architecture
 - Data Model
- features
 - Partition
 - Replication
- advantage
- existing problems

Cassandra (very rough example, more details will be concluded) Cassandra is an open source DBMS, with a peer-to-peer distributed system across its nodes, in this way it achieves a high availability. Its data model is inherited from BigTable, with the keyspaces being the container of data, in which contains column families. Row-keys are the identifiers of columns. Cassandra could always achieve a better performance by querying on row-keys.

As a distributed database management system, it partitions the data across all available nodes in the cluster. There are two partitioning strategies in Cassandra: Random partitioning (default and recommended) and Ordered partitioning. If partitioning is initialized, every change made will not be processed without reloading the whole data again, so it can be quite computationally expensive.

In every transaction, Cassandra replicates data to achieve high availability by storing replicas on multiple nodes in the cluster to guarantee reasonable fault tolerance. So Cassandra rely on a high speed and low latency network connection. Since it does not operate under the master-slave replication strategy like in HBase, thus situations like the single points failures do not exist in

Cassandra, ensuring higher level of availability. But in this case, the extra transactions incurred between could be expensive.

In the CAP(Consistency, Availability, Partition tolerance) theorem, Cassandra satisfied the Availability and Partition tolerance, these are its major advantages. It needs to trade-off between availability and consistency depending on the actual context of tasks.