

# Lab 5

Jiarong Ye

September 20, 2018

## import packages

```
In [34]: import datascience as ds
        from datascience import *
        import numpy as np
        from graphviz import Source
        import pandas as pd
        import re
        import string
        from copy import copy
        import nltk
        from datetime import datetime

        from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
        from sklearn.model_selection import train_test_split
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import RandomForestClassifier
        from xgboost import XGBClassifier as XGBoostClassifier
        from sklearn import tree
        from sklearn import metrics
        from sklearn.metrics import confusion_matrix
        import matplotlib.pyplot as plt
        from sklearn.model_selection import cross_val_score
        import seaborn as sns
        from imblearn.over_sampling import SMOTE

        %matplotlib inline

In [35]: tagged = pd.read_csv('cleaned_tagged.csv', sep=',')
        Relevant = np.array([0]*tagged.size)
        Relevant[tagged.sentiment!=-1].index.tolist() = 1
        tagged['Relevant'] = pd.Series(Relevant, name='Relevant')
        tagged = tagged.drop(tagged.columns[0], axis=1).reset_index(drop=True)
        tagged.to_csv('relevant_binary_tagged.csv')
```

## 1) tweets data loaded into Jupyter Notebook as Table object

```
In [36]: df = ds.Table.read_table('relevant_binary_tagged.csv', sep=',')
df
```

```
Out[36]: Unnamed: 0 | user_id          | user_name          | tweet_time
0                | 802657195661742080 | Christine Warren   | Wed Sep 12 01:38:14
1                | 1039245812230893570 | Trumpservative     | Wed Sep 12 01:38:16
2                | 282084840          | Darrel Sheldon #MAGAVETERAN | Wed Sep 12 01:38:18
3                | 62315639           | Queer Liberal Voting Snowflake | Wed Sep 12 01:38:18
4                | 823307049266245633 | don jones #veteran (K) | Wed Sep 12 01:38:19
5                | 340428574          | DelcoGal           | Wed Sep 12 01:38:21
6                | 1603928228         | Julz               | Wed Sep 12 01:38:22
7                | 1865678516         | Barbara Kuczinski  | Wed Sep 12 01:38:22
8                | 59288409           | Josh Steed PhD     | Wed Sep 12 01:38:25
9                | 325172419          | Mrs. Linz          | Wed Sep 12 01:38:25
... (1286 rows omitted)
```

## data preprocessing

```
In [37]: relevant_tweet = df.where('Relevant', are.equal_to(1))
irrelevant_tweet = df.where('Relevant', are.equal_to(0))
```

```
In [38]: relevant_tweet_cnt = relevant_tweet.num_rows
irrelevant_tweet_cnt = irrelevant_tweet.num_rows
```

```
In [39]: relevant_tweet_cnt, irrelevant_tweet_cnt
```

```
Out[39]: (1133, 163)
```

```
In [40]: relevant_training_size = round(relevant_tweet_cnt*0.8)
irrelevant_training_size = round(irrelevant_tweet_cnt*0.8)
```

```
In [41]: relevant_training_size, irrelevant_training_size
```

```
Out[41]: (906, 130)
```

```
In [42]: relevant_train_idx = list(range(relevant_training_size))
irrelevant_train_idx = list(range(irrelevant_training_size))
relevant_test_idx = list(range(relevant_tweet_cnt))[int(relevant_training_size):]
irrelevant_test_idx = list(range(irrelevant_tweet_cnt))[irrelevant_training_size:]
```

```
relevant_training_text = list(relevant_tweet.take(relevant_train_idx)['text'])
irrelevant_training_text = list(irrelevant_tweet.take(irrelevant_train_idx)['text'])
relevant_training_relevant = list(relevant_tweet.take(relevant_train_idx)['Relevant'])
irrelevant_training_relevant = list(irrelevant_tweet.take(irrelevant_train_idx)['Relevant'])
```

```
relevant_testing_text = list(relevant_tweet.take(relevant_test_idx)['text'])
irrelevant_testing_text = list(irrelevant_tweet.take(irrelevant_test_idx)['text'])
```

```

relevant_testing_relevant = list(relevant_tweet.take(relevant_test_idx)['Relevant'])
irrelevant_testing_relevant = list(irrelevant_tweet.take(irrelevant_test_idx)['Relevant'])

X_train = relevant_training_text + irrelevant_training_text
y_train = relevant_training_relevant + irrelevant_training_relevant
X_test = relevant_testing_text + irrelevant_testing_text
y_test = relevant_testing_relevant + irrelevant_testing_relevant

```

**2) showing the sizes of your training and testing sets (number of relevant training + irrelevant training, number of relevant testing + irrelevant testing)**

```
In [43]: len(X_train), len(X_test), len(y_train), len(y_test)
```

```
Out[43]: (1036, 260, 1036, 260)
```

**Model (Bag of Words + DTree)**

```

In [44]: vect = CountVectorizer(
            analyzer="word", ngram_range=(1,2), tokenizer=nlk.word_tokenize,
            preprocessor=None, stop_words='english', max_features=3000)
# vect = TfidfVectorizer(sublinear_tf=True, min_df=10, norm='l1', encoding='latin-1',
#                         ngram_range=(1,2), stop_words='english')
X_train_vect = vect.fit_transform(X_train).todense()
X_test_vect = vect.transform(X_test)
X_train_vect.shape, X_test_vect.shape

```

```
Out[44]: ((1036, 3000), (260, 3000))
```

```

In [45]: clf = DecisionTreeClassifier(criterion = 'entropy',
                                     random_state = 100,
                                     max_depth = 5,
                                     min_samples_leaf = 2)
# clf = RandomForestClassifier(random_state=100,
#                             n_estimators=60,
#                             criterion='entropy',
#                             n_jobs=4)
# clf = XGBoostClassifier(max_depth=5, n_estimators=5)
# clf = LogisticRegression()
clf.fit(X_train_vect, y_train)

```

```

Out[45]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=5,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=2, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=100,
                                splitter='best')

```

```

In [33]: y_pred = clf.predict(X_test_vect)
          np.mean([cross_val_score(clf, X_train_vect, y_train, cv=8)])

```

```
Out[33]: 0.8745339960944435
```

3) the confusion matrix of evaluating your Decision Tree-based relevant classifier using TESTING data, and identify false positive and false negative in your evaluation result.

```
In [46]: cm = confusion_matrix(y_test, y_pred)
         cm
```

```
Out[46]: array([[ 0, 33],
               [ 0, 227]])
```

```
In [47]: FP = cm[0][1]
         FN = cm[1][0]
         TP = cm[1][1]
         TN = cm[0][0]
         print('The False Positive:', FP)
         print('The False Negative:', FN)
         print('The True Positive:', TP)
         print('The True Negative:', TN)
```

The False Positive: 33

The False Negative: 0

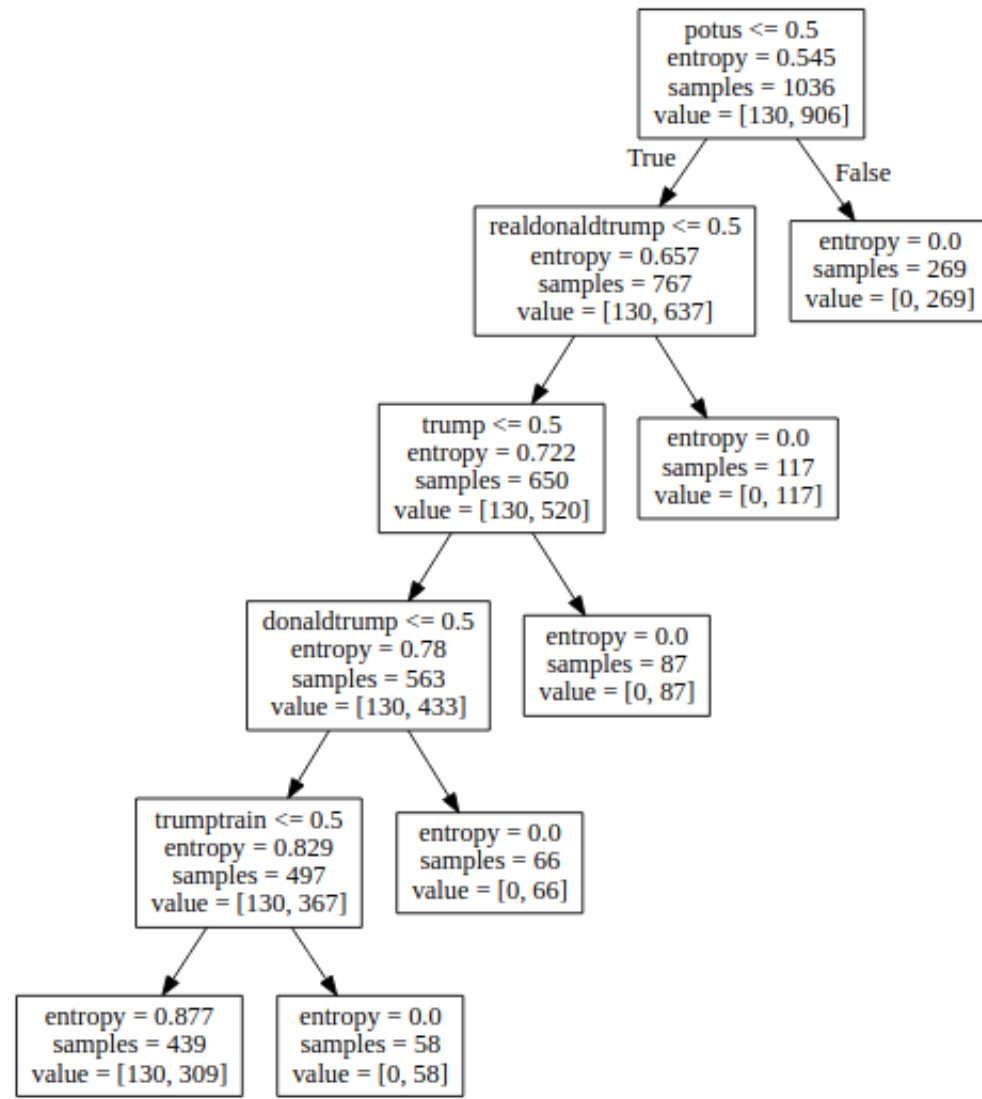
The True Positive: 227

The True Negative: 0

```
In [53]: dot_data = tree.export_graphviz(clf, out_file=None, feature_names=vect.get_feature_names()
         graph = Source(dot_data)
         graph.render('RelevantClassifier')
```

```
Out[53]: 'RelevantClassifier.pdf'
```

4) A visualization of your decision tree.



5) A description of a rule based on the tree

- If  $\text{potus} \leq 0.5$ :
  - If  $\text{realdonaldtrump} \leq 0.5$ :
    - \* If  $\text{trump} > 0.5$ :
      - then the tree reaches the leaf,  $\text{value} = [0, 87]$ , indicating that there are 87 tweets are categorized as Relevant

6) A short summary of your evaluation result using the confusion matrix

```
In [58]: def plot_cm(cm, title='Confusion matrix', cmap=plt.cm.Blues, classes=['Relevant', 'Irrelevant']):
         print(cm)
```

```

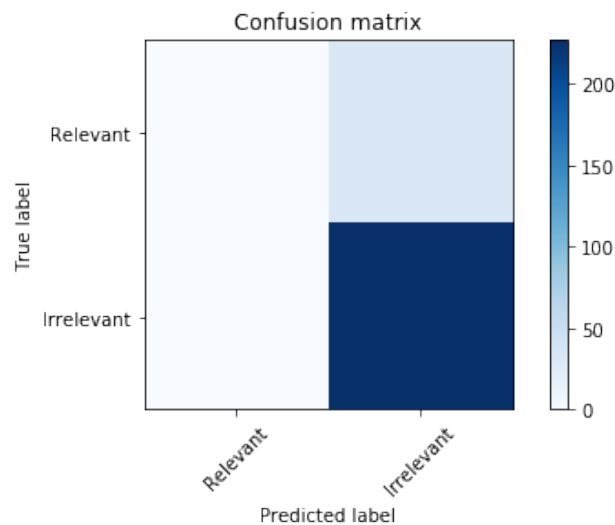
plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
target_names = classes
plt.xticks(tick_marks, target_names, rotation=45)
plt.yticks(tick_marks, target_names)
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plot_cm(confusion_matrix(y_test, y_pred))

```

```

[[ 0 33]
 [ 0 227]]

```



```

In [56]: print("AUC(ROC): " + str(metrics.roc_auc_score(y_test, y_pred)))
          print("Precision: " + str(metrics.precision_score(y_test, y_pred)))
          print("Recall: " + str(metrics.recall_score(y_test, y_pred)))
          print("F1 score: " + str(metrics.f1_score(y_test, y_pred)))

```

```

AUC(ROC): 0.5
Precision: 0.8730769230769231
Recall: 1.0
F1 score: 0.5322381930184805

```

Here we can see from the confusion matrix that the True Negative=0, and False Positive=33, which means all 33 Irrelevant tweets are categorized as Relevant. Thus there appears to have a serious problem here: even with the accuracy is decent, it is actually inflated.

The reason for that is because the proportion of Relevant tweets(1133/1296) is way larger than that of Irrelevant tweets(163/1296), thus

$$accuracy = \frac{TP + TN}{TP + TN + FN + FP} = \frac{227}{260} = 87.3\%$$

In fact, with the task being to filter out the irrelevant tweets, this classifier basically fails the task because of the unbalanced data.

### Possible fixes:

- Resampling
  - Downsampling on the Relevant tweets
  - Upsampling on the Irrelevant tweets

```
In [61]: def classify(X_train_vect, y_train, X_test_vect, y_test, model, vector):
    clf.fit(X_train_vect, y_train)
    y_pred = clf.predict(X_test_vect)
    cross_val_s = np.mean([cross_val_score(model, X_train_vect, y_train, cv=8)])
    print('Cross_val_score: {}'.format(cross_val_s))
    print("AUC(ROC): " + str(metrics.roc_auc_score(y_test, y_pred)))
    print("Precision: " + str(metrics.precision_score(y_test, y_pred)))
    print("Recall: " + str(metrics.recall_score(y_test, y_pred)))
    print("F1 score: " + str(metrics.f1_score(y_test, y_pred)))
    print(metrics.confusion_matrix(y_test, y_pred))
    plt.figure(figsize=(5,5))
    plot_cm(metrics.confusion_matrix(y_test, y_pred))
    plt.show()
    return cross_val_s
```

### Experiment 1: Downsample the Relevant data

```
In [62]: def undersample(data, percentage, vector):
    relevant_idx = tagged[tagged.Relevant==1].index
    irrelevant_idx = tagged[tagged.Relevant==0].index
    undersample_relevant_idx = np.array(np.random.choice(relevant_idx, (int(percentage*len(relevant_idx))))
    undersample_idx= np.concatenate([irrelevant_idx, undersample_relevant_idx])
    undersample_data = tagged.iloc[undersample_idx,:]
    print('There are {} relevant tweets.\n There are {} irrelevant tweets.'.format(len(relevant_idx), len(irrelevant_idx)))

    X_train, X_test, y_train, y_test = train_test_split(undersample_data.text, undersample_data.tagged, test_size=0.2, random_state=42)
    X_train_vect = vector.fit_transform(X_train).todense()
    X_test_vect = vector.transform(X_test)
    plt.figure(figsize=(5,5))
    sns.countplot("Relevant", data=undersample_data)
    plt.show()
    return X_train_vect, X_test_vect, y_train, y_test
```


```

In [66]: cross_val_scores=[]
         for percentage in np.arange(1, 4, 0.5):
             X_train_vect, X_test_vect, y_train, y_test = undersample(tagged, percentage, vect)
             cross_val_s = classify(X_train_vect=X_train_vect,
                                   X_test_vect=X_test_vect,
                                   y_train =y_train,
                                   y_test=y_test,
                                   model=clf,
                                   vector=vect)
             cross_val_scores.append(cross_val_s)
         plt.figure(figsize=(5,5))
         plt.plot(np.arange(1,4,0.5), cross_val_scores)
         plt.scatter(np.arange(1,4,0.5), cross_val_scores, c='r')
         plt.show()

```

There are 163 relevant tweets.

There are 163 irrelevant tweets.



Cross\_val\_score: 0.8041758308895406

AUC(ROC): 0.8157894736842105

Precision: 1.0

Recall: 0.631578947368421

F1 score: 0.7741935483870968


[[28 0]

[14 24]]


[[28 0]

[14 24]]






There are 244 relevant tweets.  
There are 163 irrelevant tweets.




Cross\_val\_score: 0.824130315822389  
AUC(ROC): 0.8773584905660378  
Precision: 1.0

Recall: 0.7547169811320755  
F1 score: 0.8602150537634409  
[[29 0]  
[13 40]]  
[[29 0]  
[13 40]]



There are 326 relevant tweets.  
There are 163 irrelevant tweets.



Cross\_val\_score: 0.7495153061224489

AUC(ROC): 0.875

Precision: 1.0

Recall: 0.75


F1 score: 0.8571428571428571

```
[[30  0]
```

```
 [17 51]]
```


```
[[30  0]
```

```
 [17 51]]
```



There are 407 relevant tweets.

There are 163 irrelevant tweets.



Cross\_val\_score: 0.7327586206896551

AUC(ROC): 0.8192771084337349

Precision: 1.0

Recall: 0.6385542168674698


F1 score: 0.7794117647058824

[[31 0]


[30 53]]

[[31 0]

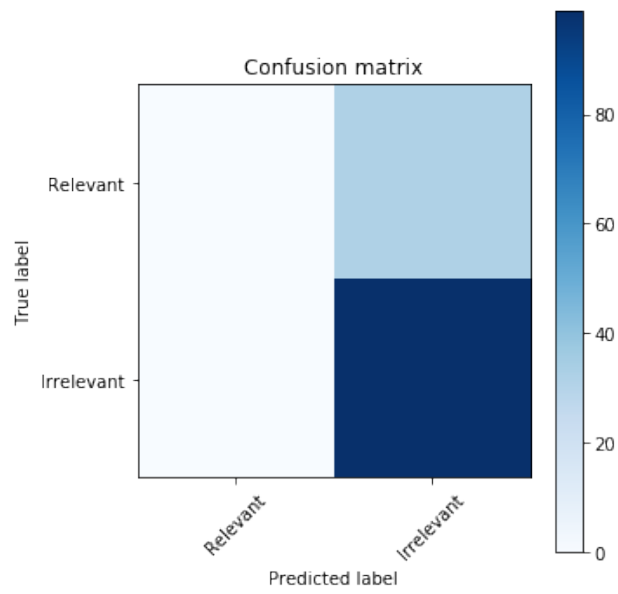
[30 53]]




There are 489 relevant tweets.  
There are 163 irrelevant tweets.



Cross\_val\_score: 0.7214352054195804  
AUC(ROC): 0.5  
Precision: 0.7557251908396947  
Recall: 1.0  
F1 score: 0.8608695652173913  
[[ 0 32]  
 [ 0 99]]  
[[ 0 32]  
 [ 0 99]]


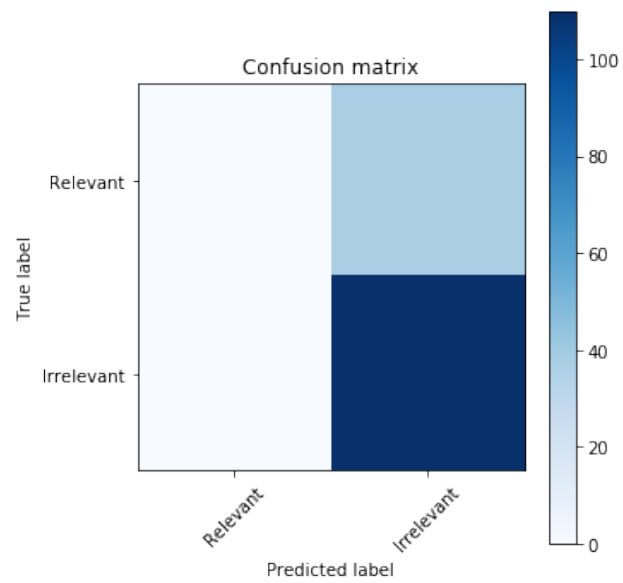


There are 570 relevant tweets.  
There are 163 irrelevant tweets.



Cross\_val\_score: 0.7850140380106133  
AUC(ROC): 0.5  
Precision: 0.7482993197278912

Recall: 1.0  
F1 score: 0.8560311284046693  
[[ 0 37]  
[ 0 110]]  
[[ 0 37]  
[ 0 110]]



## Experiment 2: Upsample the Irrelevant data

```
In [64]: def oversample(data, vector):  
        os = SMOTE(random_state=0)  
        X_train, X_test, y_train, y_test = train_test_split(data.text, data.Relevant, test_  
        X_train_vect = vector.fit_transform(X_train).todense()  
        X_test_vect = vector.transform(X_test)  
        os_X_train, os_y_train = os.fit_sample(X_train_vect, y_train)  
        return os_X_train, X_test_vect, os_y_train, y_test
```

```
In [65]: X_train_vect, X_test_vect, y_train, y_test = oversample(tagged, vect)  
        classify(X_train_vect=X_train_vect,  
                X_test_vect=X_test_vect,  
                y_train =y_train,  
                y_test=y_test,  
                model=clf,  
                vector=vect)
```

Cross\_val\_score: 0.8281613879832324

AUC(ROC): 0.8223684210526316

Precision: 1.0

Recall: 0.6447368421052632

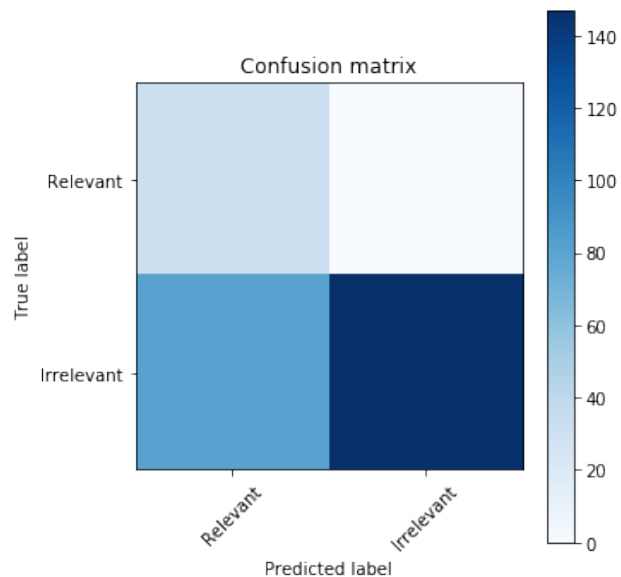
F1 score: 0.7839999999999999

```
[[ 32  0]
```

```
 [ 81 147]]
```

```
[[ 32  0]
```

```
 [ 81 147]]
```



Out [65]: 0.8281613879832324